


API Gateway for Docker Containers

Problem statement

Design an API Gateway CLI tool, *rgate*



rgate can be run with a command like below to receive traffic on localhost:8080

```
rgate --port 8080 --config config.yml
```

With the contents of *config.yml* looks like,

```
routes:
  - path_prefix: /api/payment
    backend: payment
  - path_prefix: /api/orders
    backend: orders
default_response:
  body: "This is not reachable"
  status_code: 403
```

```

backends:
  - name: payment
    match_labels:
      - app_name=payment
      - env=production
  - name: orders
    match_labels:
      - app_name=orders
      - env=production

```

Features

- Each backend is a container.
- A container is selected as a backend if the container has all the *match_labels* present as Docker labels. If there are multiple containers matching, select a random container as a backend.
- Incoming http requests are routed to the corresponding backend if the path starts with the given *path_prefix*.
- If there are no routes matching the request, it should respond with the given body and *status_code* in the *default_response*
- If the backend is down, respond with 503 code
- Accessing `http://localhost:8080/stats` should give us information about the traffic it has received so far. An example JSON response is

```

{
  "requests_count" : {
    "success": 100, // status codes 200-399
    "error": 110,   // status codes >399
  },
  "latency_ms": {
    "average": 2,
    "p95": 5,
    "p99": 10
  }
}

```

Instructions

- Use any language that you are comfortable with
- You are NOT expected to implement any backend yourself. You could use an nginx container in your local to test
- Aim to achieve good unit test coverage
- Ensure README has instructions for installing the tool, testing, usage, etc.. Feel free to add your thoughts on the design choices and assumptions